

COREDNS INTRO
KUBECON, COPENHAGEN
MIEK GIEBEN
MAY 2-4, 2018

In this presentation

- CoreDNS status and roadmap-ish
- CoreDNS in k8s
- Debugging features
- Plugins
 - Write your own: NXDOMAIN domains

STATUS

CoreDNS: flexible DNS server powered by plugins
(support DNS, DNS over TLS, DNS over gRPC)

- Picking new developers (thanks!)
- Moved to OWNERS system for plugins
- Security review from Cure53 (sponsored by the CNCF)
 - one big hole plugged with 1.1.1 (cache spoofing)
 - 2 other minor bugs
- Release cadence: every 3 to 4 weeks

ROADMAP(-ISH)

- Core (non-plugin side) very stable; occasional feature enhancements
- Plugins
 - *kubernetes*: small tweaks here and there
 - *etcd3* in the works
 - *sql* serve from database (in the works)
 - note: *explugins/pdsql*
- **Watch** functionality:
 - gRPC based, register interested in *name*
 - get ping when name changes

COREDNS IN K8S

GA in 1.11 (I hope) some issues need to be worked out.

- pushing gcr.io - CoreDNS being the guinea pig for a new process

"This is progressing"

COREDNS CONFIGURATION IN K8S

Kubernetes setup comes with this Corefile:

```
.:53 {  
  errors  
  health  
  kubernetes cluster.local 10/8 {  
    pods insecure  
    upstream  
    fallthrough in-addr.arpa  
  }  
  prometheus :9153  
  proxy . /etc/resolv.conf  
  cache 30  
}
```

Each line a plugin.

. : 53 a server for *everything* on port 53.

ERRORS

Log errors from plugins to standard output.

i.e. `unreachable backend: <error>` from the *proxy* plugin.

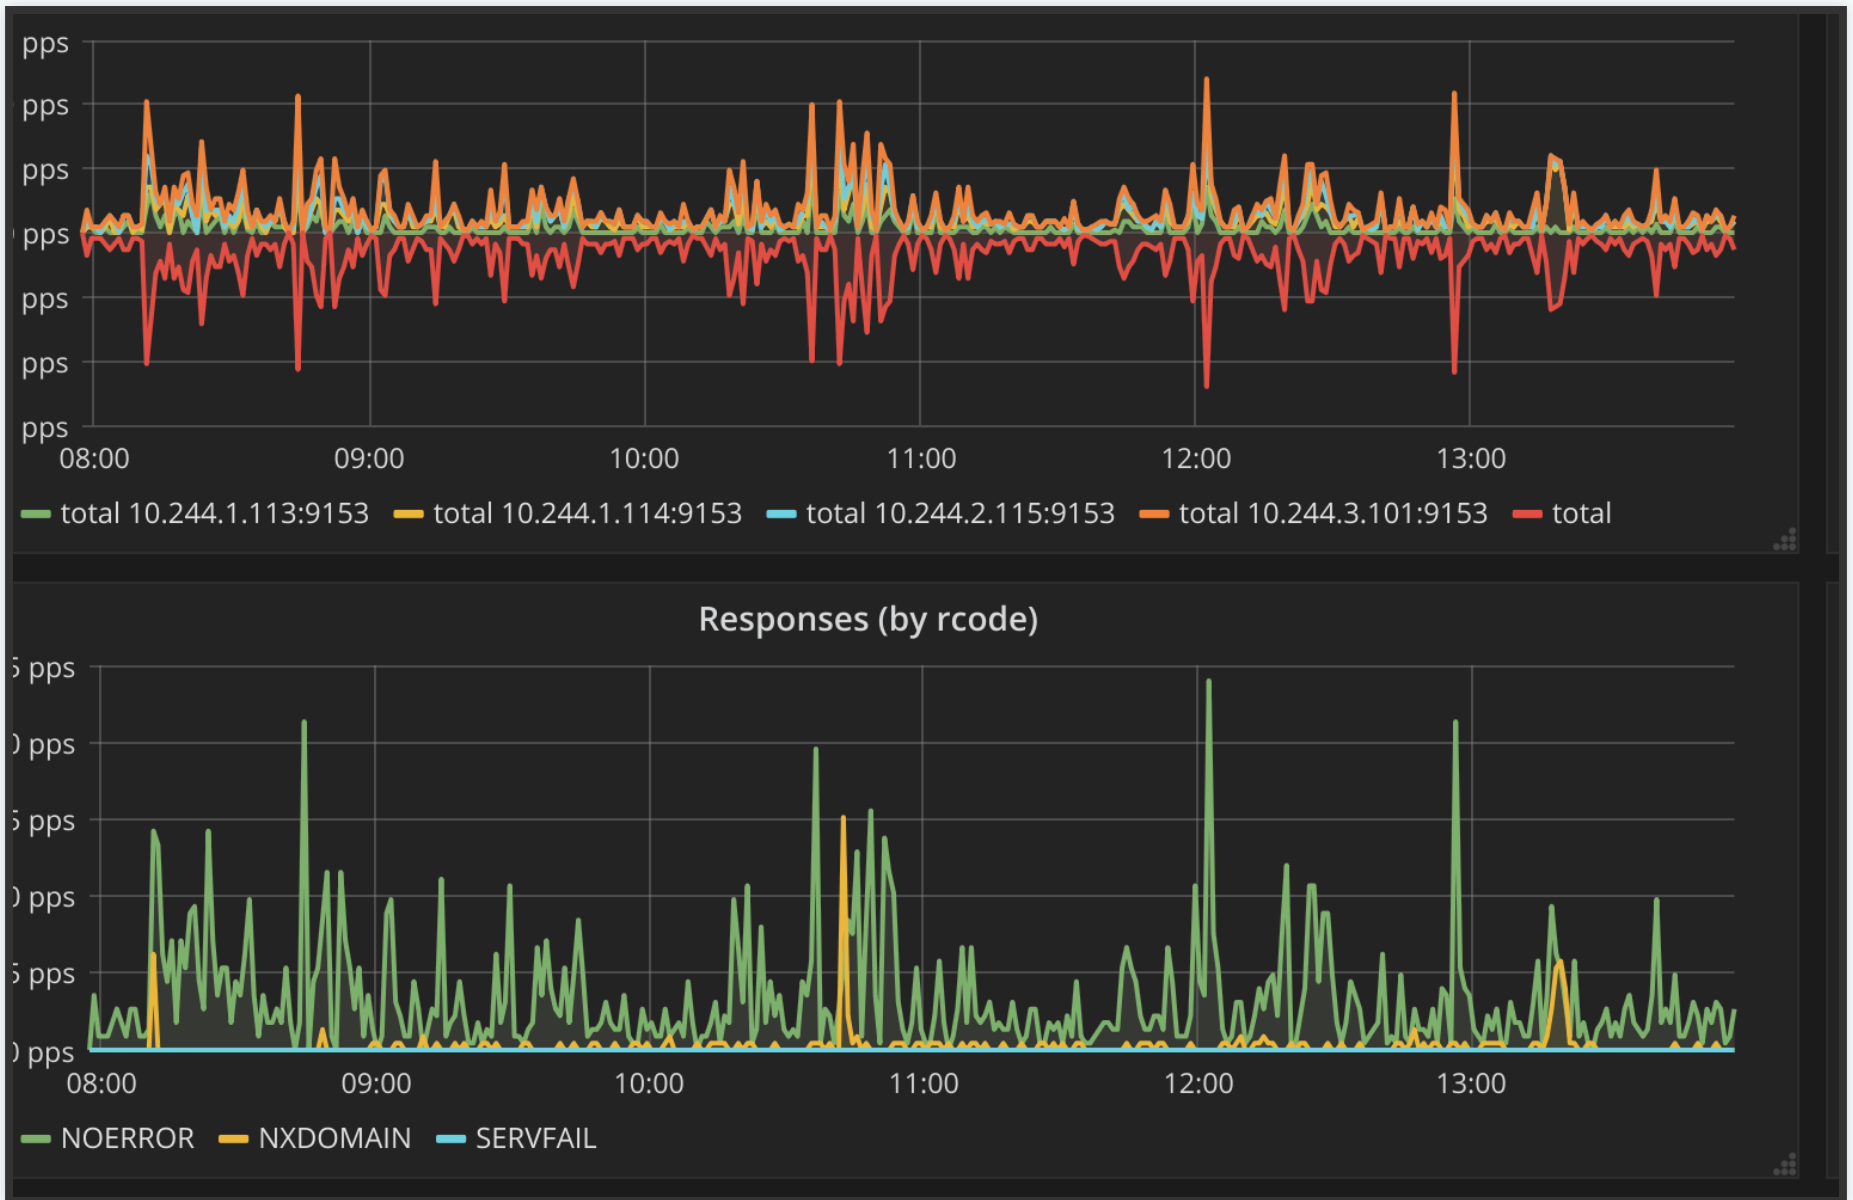
Note: all plugins have documentation at:
<https://coredns.io/plugins/<NAME>>

HEALTH

Expose health on port 8080. OK once we've synced to k8s api. (SERVFAIL until that happens)

PROMETHEUS

Metrics exposed on 9153; queries, latency, etc.



Grafana dashboard

PROXY

Forward queries *not* for `cluster.local` to
nameserver(s) defined in `/etc/resolv.conf`

CACHE

Cache every (sane) response up to 30 seconds. Can do other stuff: *prefetching*, separate *positive/negative* cache.

KUBERNETES

Kubernetes provides K8S service discovery.

```
kubernetes cluster.local 10/8 {  
  pods insecure  
  upstream  
  fallthrough in-addr.arpa  
}
```

- handle cluster.local
- 10/8: 10.in-addr.arpa (10.x.x.x reverse lookups)

DIRECTIVES

- `pod` `insecure`: always return A record for IP requests
`1-2-3-4.default.pod.cluster.local.` → `1.2.3.4.`
- `upstream`: resolve external names.
- `fallthrough` `in-addr.arpa` - reverse (PTR) lookups, fallthrough to proxy when NXDOMAIN.

NEED MORE?

log and maybe debug?

- log: log all processed queries to standard output
- debug: enable debug logging
 - disables recover after panic

DEBUGGING

- with *debug* log. Debug enabled output shown.
- zone transfers in *kubernetes*: Give me **all** records.
enable with "transfer to *", check with
dig AXFR cluster.local.
- *health* exports metric:
coredns_health_request_duration_seconds
- (external) plugin *dump*: log everything *before* any
processing. (with tcpdump).

feedback/ideas/help welcome!

PROBLEM

dnsmasq has this feature: NXDOMAIN domains

```
server=/example.net/
```

Shortcuts the resolving of this domain.

I want this in CoreDNS and don't want to deepdive into *rewrite* or *template*.

Let's write a plugin **that does the same thing**.
(github.com/miekg/nxdomain)

PLUGIN

For a plugin to work, we need:

1. (Go) code that implements the `plugin.Handler` interface:
2. A setup function that parses the Corefile.
3. To register it into CoreDNS.

HANDLER INTERFACE

```
Name() string
ServeDNS(context.Context, dns.ResponseWriter, *dns.Msg) \
(int, error)
```

- **Name** return name of the plugin (nxdomain)
- **ServerDNS** guts of the plugin: query handling

```
// N implements the plugin interface.
type N struct {
    Next plugin.Handler // needed for chaining the plugins
    names []string       // domains we shortcut
}
```

SETUP

```
func setup(c *caddy.Controller) error {  
    // Parse: nxdomain <list of domains>  
    names := []string{}  
    for c.Next() {  
        args := c.RemainingArgs()  
        if len(args) == 0 {  
            return plugin.Error("nxdomain", c.ArgErr())  
        }  
        // I'll bet these are not fully qualified.  
        for _, a := range args {  
            names = append(names, dns.Fqdn(a))  
        }  
    }  
    // ...  
}
```

REGISTER IN COREDNS

```
// ...
dnsserver.GetConfig(c).AddPlugin(func(
    next plugin.Handler) plugin.Handler {
    return N{Next: next, names: names}
})

return nil
}
```

SERVEDNS

```
func (n N) ServeDNS(ctx context.Context,  
    w dns.ResponseWriter, r *dns.Msg) (int, error) {  
    for _, n := range n.names {  
        if dns.IsSubDomain(n, r.Question[0].Name) {  
            m := new(dns.Msg)  
            m.SetRcode(r, dns.RcodeNameError)  
            m.Ns = []dns.RR{soa(n)}  
            w.WriteMsg(m)  
            return 0, nil  
        }  
    }  
    return plugin.NextOrFailure(n.Name(), n.Next, ctx, w, r)  
}
```

COMPILING COREDNS

In `plugin.cfg`

```
...  
debug:debug  
trace:trace  
health:health  
nxdomain:github.com/miekg/nxdomain  
pprof:pprof  
prometheus:metrics  
errors:errors  
log:log  
...
```

\$ go generate && go build

Check if it's there: \$ coredns -plugins (should list `dns.nxdomain`)

TRYING OUT

```
. {  
  nxdomain example.org  
  whoami  
}
```

```
$ coredns -dns.port=1043  
$ dig @localhost -p 1043 example.net  
;; ADDITIONAL SECTION:  
example.net.      0    IN    AAAA    ::1  
_udp.example.net. 0    IN    SRV 0 0 38755 .  
  
$ dig @localhost -p 1053 example.org  
->>HEADER<<- opcode: QUERY, status: NXDOMAIN, id: 46392
```

That's all! Questions?